

CSE455 Machine Learning

Homework 3

Sedef Erdoğan, 1801042102

May 13, 2023

0.1 Import Necessary Libraries

```
[1]: from keras.datasets import mnist
import numpy as np
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.decomposition import KernelPCA
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score
```

```
2023-05-13 22:55:55.963536: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not
find cuda drivers on your machine, GPU will not be used.
2023-05-13 22:55:56.039377: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not
find cuda drivers on your machine, GPU will not be used.
2023-05-13 22:55:56.040642: I tensorflow/core/platform/cpu_feature_guard.cc:182]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.
2023-05-13 22:55:58.561717: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not
find TensorRT
```

0.2 Read Dataset

Load MNIST Dataset

```
[2]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

0.3 PCA Implementation

```
[3]: def pca(X):
    mean = np.mean(X, axis=0)

    X_normalized = X - mean

    U, S, Vt = np.linalg.svd(X_normalized)
```

```

weights = S ** 2 / (X.shape[0] - 1)

vectors = Vt.T

return mean, weights, vectors

```

0.4 Using PCA before Classification

```
[4]: x_train.shape
```

```
[4]: (60000, 28, 28)
```

```
[5]: x_test.shape
```

```
[5]: (10000, 28, 28)
```

```

[6]: X_train = x_train.reshape(x_train.shape[0], -1)
      X_test = x_test.reshape(x_test.shape[0], -1)

      print('Train Shape:', X_train.shape)
      print('Class Shape:', X_test.shape)

```

Train Shape: (60000, 784)

Class Shape: (10000, 784)

Selecting a subset of data

```

[7]: subset_size = 1000
      indices = np.random.choice(X_train.shape[0], size=subset_size, replace=False)
      x_train_subset = X_train[indices]
      y_train_subset = y_train[indices]

```

```

[8]: num_components = [3, 5, 10, 20] # Choose the number of components for reduction

      for n in num_components:
          mean, weights, vectors = pca(x_train_subset)

          reduced_X_train = np.dot(x_train_subset - mean, vectors[:, :n])
          reduced_X_test = np.dot(X_test - mean, vectors[:, :n])

          plt.figure(figsize=(8, 6))

          for i in range(10):
              indices = np.where(y_test == i)[0]
              plt.scatter(reduced_X_test[indices, 0], reduced_X_test[indices, 1], u
→label=str(i))

```

```

plt.xlabel("First Principal Component")
plt.ylabel("Second Principal Component")
plt.legend()
plt.title("MNIST Test Data - PCA with {} Components".format(n))
plt.show()

if n >= 3:
    plt.figure(figsize=(8, 6))
    for i in range(10):
        indices = np.where(y_test == i)[0]
        plt.scatter(reduced_X_test[indices, 0], reduced_X_test[indices, 2],
→label=str(i))
        plt.xlabel("First Principal Component")
        plt.ylabel("Third Principal Component")
        plt.legend()
        plt.title("MNIST Test Data - PCA with {} Components".format(n))
        plt.show()

    # Feed reduced features to Random Forest and perform classification with
→cross-validation
    rf = RandomForestClassifier(n_estimators=100)

    accuracies = []

    # Perform cross-validation
    kf = KFold(n_splits=5, shuffle=True)

    for train_index, test_index in kf.split(reduced_X_train):
        x_cv_train, x_cv_test = reduced_X_train[train_index],
→reduced_X_train[test_index]
        y_cv_train, y_cv_test = y_train_subset[train_index],
→y_train_subset[test_index]


        rf.fit(x_cv_train, y_cv_train)


        y_pred = rf.predict(x_cv_test)
        accuracy = accuracy_score(y_cv_test, y_pred)
        accuracies.append(accuracy)

    # Calculate the average accuracy over all folds
    avg_accuracy = np.mean(accuracies)


    print("Average accuracy with {} components: {:.2f}%".format(n, avg_accuracy
→* 100))


```







Average accuracy with 3 components: 44.00%







Average accuracy with 5 components: 65.20%





Average accuracy with 10 components: 81.20%





Average accuracy with 20 components: 85.40%

0.5 Non-linear Versions of PCA

```
[9]: num_components = [3, 5, 10, 20]  # Choose the number of components for reduction

for n in num_components:
    kpca = KernelPCA(n_components=n, kernel='poly', degree=3)
    reduced_X_train = kpca.fit_transform(x_train_subset)
    reduced_X_test = kpca.transform(X_test)

    plt.figure(figsize=(8, 6))

    for i in range(10):
        indices = np.where(y_test == i)[0]
        plt.scatter(reduced_X_test[indices, 0], reduced_X_test[indices, 1],
            ↪label=str(i))
    plt.xlabel("First Principal Component")
    plt.ylabel("Second Principal Component")
    plt.legend()
```

```

plt.title("MNIST Test Data - PCA with {} Components".format(n))
plt.show()

if n >= 3:
    plt.figure(figsize=(8, 6))
    for i in range(10):
        indices = np.where(y_test == i)[0]
        plt.scatter(reduced_X_test[indices, 0], reduced_X_test[indices, 2],
→label=str(i))
        plt.xlabel("First Principal Component")
        plt.ylabel("Third Principal Component")
        plt.legend()
        plt.title("MNIST Test Data - PCA with {} Components".format(n))
        plt.show()

    # Feed reduced features to Random Forest and perform classification with
→cross-validation
    rf = RandomForestClassifier(n_estimators=100)

    accuracies = []

    # Perform cross-validation
    kf = KFold(n_splits=5, shuffle=True)

    for train_index, test_index in kf.split(reduced_X_train):

        x_cv_train, x_cv_test = reduced_X_train[train_index],
→reduced_X_train[test_index]
        y_cv_train, y_cv_test = y_train_subset[train_index],
→y_train_subset[test_index]


        rf.fit(x_cv_train, y_cv_train)


        y_pred = rf.predict(x_cv_test)
        accuracy = accuracy_score(y_cv_test, y_pred)
        accuracies.append(accuracy)

    # Calculate the average accuracy over all folds
    avg_accuracy = np.mean(accuracies)


    print("Average accuracy with {} components: {:.2f}%".format(n, avg_accuracy
→* 100))

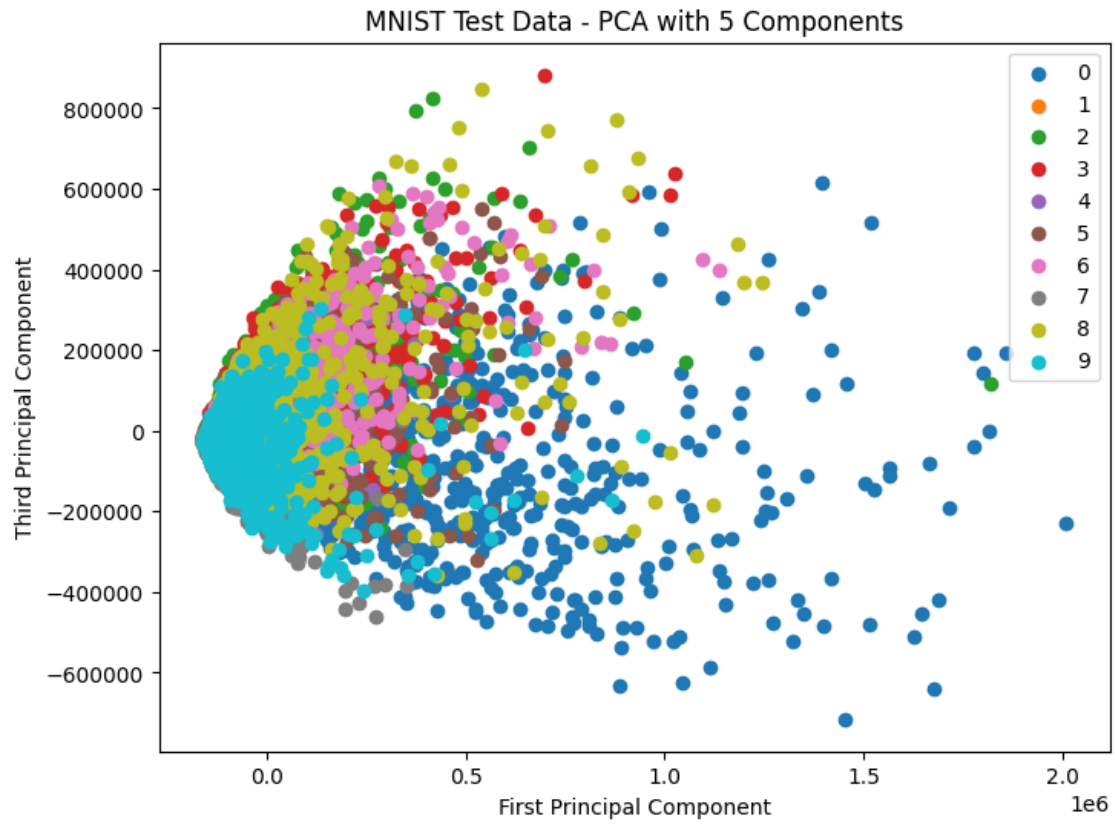
```







Average accuracy with 3 components: 37.10%







Average accuracy with 5 components: 61.10%





Average accuracy with 10 components: 77.20%





Average accuracy with 20 components: 82.50%